

Data flow analysis for Uranus applications

Chan Kwan Yin

December 17, 2020

Outline

1 Background

2 Approach


3 Limitations

4 Conclusion

Background

SGX Enclaves

- Servers outsourced to third-party cloud providers
- Threat model: Adversaries with privileged access to OS, BIOS or hardware
- Enclave protects both code and memory from these adversaries



Uranus [2]

- OpenJDK fork that supports Intel SGX
- Methods marked as `@JECall` enters enclaves until return
- Methods marked as `@JOCall` exits enclaves until return
- Useful for integration with libraries like Hadoop and Spark
- Question: Where should `@JECall` and `@JOCall` be placed?
- Question: Is code in these libraries safe as enclave code?

The problem: Performanc/Security Tradeoff

- More code outside enclave:
 - Increased risk of leaking protected data
 - Some leaks may come from unexpected side channels
- More code into enclave:
 - Limited EPC (Enclave Page Cache)
 - Up to 100 MB of EPC
 - Out of memory \implies extremely slow swap
 - JVM applications especially memory-greedy
 - Principle of Least Privilege
 - Vulnerabilities in enclave code bypass enclave protection
 - Vulnerabilities in code outside must only use specific entry points
 - Reduce attack surface

The solution

- Select the minimum cover of sensitive data
- *enclavlow*¹: an information flow analysis tool
- Sources of sensitive data marked with `sourceMarker`
- Anonymization marked with `sinkMarker`
- Identity functions; expected to be optimized them away by JIT

```
1 @JECall
2 static int process(byte[] encrypted) {
3     byte sum = 0;
4     byte[] password = sourceMarker(new byte[]{1, 2, 3, 4, 5, 6});
5     byte[] decrypted = decrypt(password, encrypted);
6     for(byte b : decrypted) sum ^= b;
7     return sinkMarker(sum);
8 }
```

¹coined from the words "enclave" and "flow"

Threat model

- Adversary has no read/write access to enclave code and memory
- Adversary has arbitrary access to any system resource, *including* non-enclave Java and JVM memory
- The actual adversary also has access to hardware resources, including sensor modules, system clock, etc.
- These resources can be used to implement side channel attacks, which are system-dependent, environment-dependent and architecture-dependent
- Example: Timing attack based on CPU branch prediction optimization
- Our adversary model excludes these side channel attacks

Changes on Uranus model

- Applications are expected to run on a Uranus-based JVM.
- Some behaviour disallowed by Uranus assumed permissible under explicit indication
- Assignment to objects outside enclaves
 - Reads are assumed copied into enclave
 - Writes are assumed always leaking
- Assignment to static fields
 - Reads are cloned into enclave
 - Writing enclave-local data may not be expected behaviour
 - Assumed immediate leak
 - Proposed `EnclaveLocal` wrapper type
- Expected to integrate into Uranus for compile-time checking/optimization

Approach

Intuition: Trivial ways of leaking data

```
1  static int outside; // static variables stored outside enclave
2  @JOCall void println(int x);
3
4  int secret() {
5      return sourceMarker(123456);
6  }
7
8  @JECall int foo() {
9      store = secret(); // assigning to static field
10     println(secret()); // passing secret to a JOCall
11     return secret(); // returning secret out of a JECall
12 }
```

Intuition: Non-trivial ways of leaking data

- Assigning to an outside-enclave object:

```
1 @JECall void x(Box box) {  
2     box.value = secret();  
3 }
```

- Leaking control flow into variables:

```
1 if(secret() >= 3) outside++;  
2 for(int i=0; i < secret(); i++) outside++;
```

- Implicit exceptions:

```
1 int[3] array = new int[3];  
2 array[secret()] = 1;
```

- Transitive application of above

Flow graph

- Analysis framework: Soot [1]
- Each readable/writable data entry as a node
- $(x, y) \in E \implies$ placement of y outside enclave reduces indistinguishability of x
- $(x, y), (y, z) \in E$: transitive
- $(x, y), (x, z) \in E$: Leaking *either* y or z distinguishes x
- No way to represent requirement of *both* y and z

Procedure

- 1 Each method is analyzed independently as a Local Flow Graph (LFG)
- 2 LFG contracted into subgraph of only "public" nodes called Contract Flow Graph (CFG)
- 3 CFGs merged together into Aggregate Flow Graph (AFG),

Local flow graph

- Analyze each method independently using Soot's `ForwardBranchedFlowAnalysis`
- Soot calls the `flowThrough` method for each 3AC (jimple) statement
- Each `flowThrough` maps the state in the previous step to a new state
- For branching statements, each branch has a its own state
- When flows converge, soot calls the `merge` method to map two states into a new one

Implementing `flowThrough` for assignment operations

- For each type of value, define
 - rvalue nodes
 - lvalue nodes for assignment
 - and lvalue nodes for deletion

- Pseudocode:

```

1  Algorithm assign($left = $right):
2    for each lvalues($left, FLAG_DELETION) as $node:
3      delete (*, $node) from flow graph
4    for each lvalues($left) as $leftNode:
5      add (CTRL, $leftNode) to flow graph
6      for each rvalues($right) as $rightNode:
7        add ($rightNode, $leftNode) to flow graph // intuitive
8    for each lvalues($right) as $leftNode:
9      for each rvalues($left) as $rightNode:
10       add ($rightNode.*, $leftNode.*) to flow graph // field projection

```


Global nodes

- Static node
 - All reads use Uranus cache
 - Writes are assumed immediate leak
- Explicit source/sink
 - Explicit sink is cosmetic

Method signature nodes

- Parameter nodes
 - Supports both read and write
- This node
 - just an alternative parameter
- Return node
- Throw node
 - just an alternative return path

The CTRL node

- Tracks what data affect the current flow
- Each branched block has its CTRL node
- $(x, \text{CTRL}) \in E \implies \exists$ predicate p such that the branch represented by CTRL executes if and only if $p(x)$
- $(\text{CTRL}_1, \text{CTRL}_2) \in E$ when CTRL_2 is a subbranch of CTRL_1
- $(\text{CTRL}, y) \in E \implies$ leaking y may distinguish whether CTRL was run
- Merging states: select Lowest Common Ancestor

Other nodes

- Local variables
 - Including temp values created by Jimple
- Method calls
 - Blackbox proxies to signature nodes

Case study: Branching (tableswitch)

```

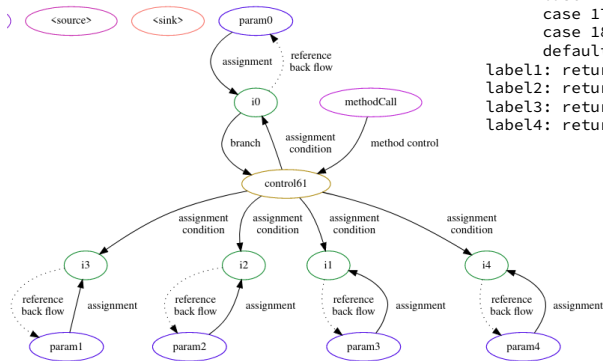
1  public static int switchMux(int x,
2      int a, int b, int c, int d) {
3      switch (x) {
4          case 16: return a;
5          case 17: return b;
6          case 18: return c;
7          default: return d;
8      }
9  }

```

```

public static int switchMux
    (int, int, int, int, int) {
    int i0, i1, i2, i3, i4;
    i0 := @parameter0: int;
    i3 := @parameter1: int;
    i2 := @parameter2: int;
    i1 := @parameter3: int;
    i4 := @parameter4: int;
    tableswitch(i0) {
        case 16: goto label1;
        case 17: goto label2;
        case 18: goto label3;
        default: goto label4; };
    label1: return i3;
    label2: return i2;
    label3: return i1;
    label4: return i4; }

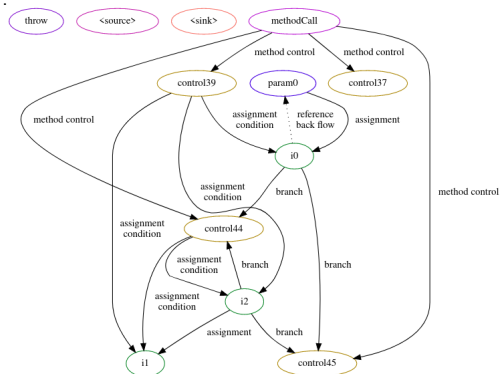
```



Case study: Loops

```
1 public static int loopInc(int i) {  
2     int a = 0;  
3     for (int j = 0; j < i; j++) {  
4         a += j;  
5     }  
6     return a;  
7 }
```

```
1 public static int loopInc(int) {  
2     int i0, i1, i2;  
3     i0 := @parameter0: int;  
4     i1 = 0;  
5     i2 = 0;  
6     label1:  
7     if i2 >= i0 goto label2;  
8     i1 = i1 + i2;  
9     i2 = i2 + 1;  
10    goto label1;  
11    label2: return i1; }
```

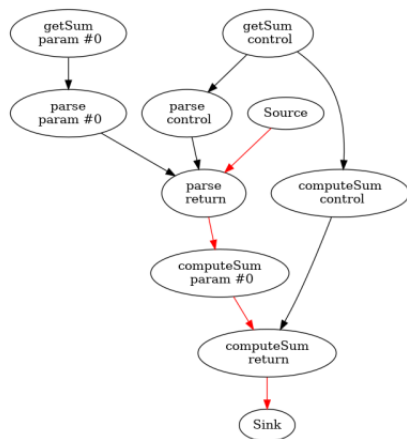


Contraction and aggregation

- At each return point, track flow from "public" nodes into the following nodes:
 - `this`
 - `static`
 - explicit sink
 - parameters
 - method calls (proxy nodes)
- At aggregation, connect each proxy node to the actual implementation
 - Polymorphism?

AFG: Full Example

```
1 @JECall
2 int getSum(byte[] enc) {
3     byte[] dec = decrypt(enc);
4     return sinkMarker(computeSum(dec)); }
5
6 byte[] decrypt(byte[] enc) {
7     byte[] otp = sourceMarker(PRIVATE_KEY);
8     byte[] result = new byte[enc.length];
9     for(int i = 0; i < buf.length; i++)
10         result[i] = enc[i] ^ otp[i];
11     return result; }
12
13 int computeSum(byte[] bytes) {
14     int result = 0;
15     for(byte b : bytes) result += (int) result;
16     return result; }
```



Limitations

False negatives: implicit exceptions

- Conditional runtime errors leaks data
 - `NullPointerException`
 - `IndexOutOfBoundsException`
- Too sensitive to raise an exception path for every array access and object access
- Good practice: exceptions should be caught at enclave boundary anyway!
- Solutions usually achieved at the language level, e.g. `@NonNull`, `@Size`
- Do not reinvent the wheel

False positives: identical branches

- This does not leak `secret` (assuming `doSomething*()` do not leak control flow)

```
1  boolean foo() {  
2      int secret = getSecret();  
3      if(secret > 1) {  
4          doSomething();  
5          return false;  
6      } else {  
7          doSomethingElse();  
8          return false;  
9      }  
10 }
```

- Suggested fix: factorize common code out of branches (good code style)

False positives: self-anonymization

- This does not leak `secret` (assuming `doSomethingWith(int)` does not leak)

```
1  int foo(int a) {  
2      int secret = getSecret();  
3      a += secret;  
4      doSomethingWith(a);  
5      a -= secret;  
6      return a;  
7  }
```

- Suggested fix: create a clone of `a`
- Immutability paradigms?
 - We are considering the *compiled binary* layer, so functional programming languages do not help

Tackling polymorphism

- Computing all combinations of instance classes requires exponential time.
- Solution: union of CFGs of all possible subclasses
 - Iago attacks is still possible if code is used elsewhere in enclave
 - Therefore all subclasses, not only those through call analysis, are considered

Tackling JNI calls

- Assumptions:
 - parameters are independent
 - return value is the output
 - control flow is not leaked
- False negative example: `System.arraycopy`
- Uranus turns system calls into `OCalls`, which leak control flow.

Integration into Uranus

- Uranus disallows reads/writes of objects outside enclave without `SafeGetField` etc
- Runtime overhead of checking object location
- *enclavlow* to perform this analysis at compile-time

Conclusion

- Assist with decisions on performance/security tradeoff
- Incorporated into Uranus
- Applications in big data industry
- Room for improvement

